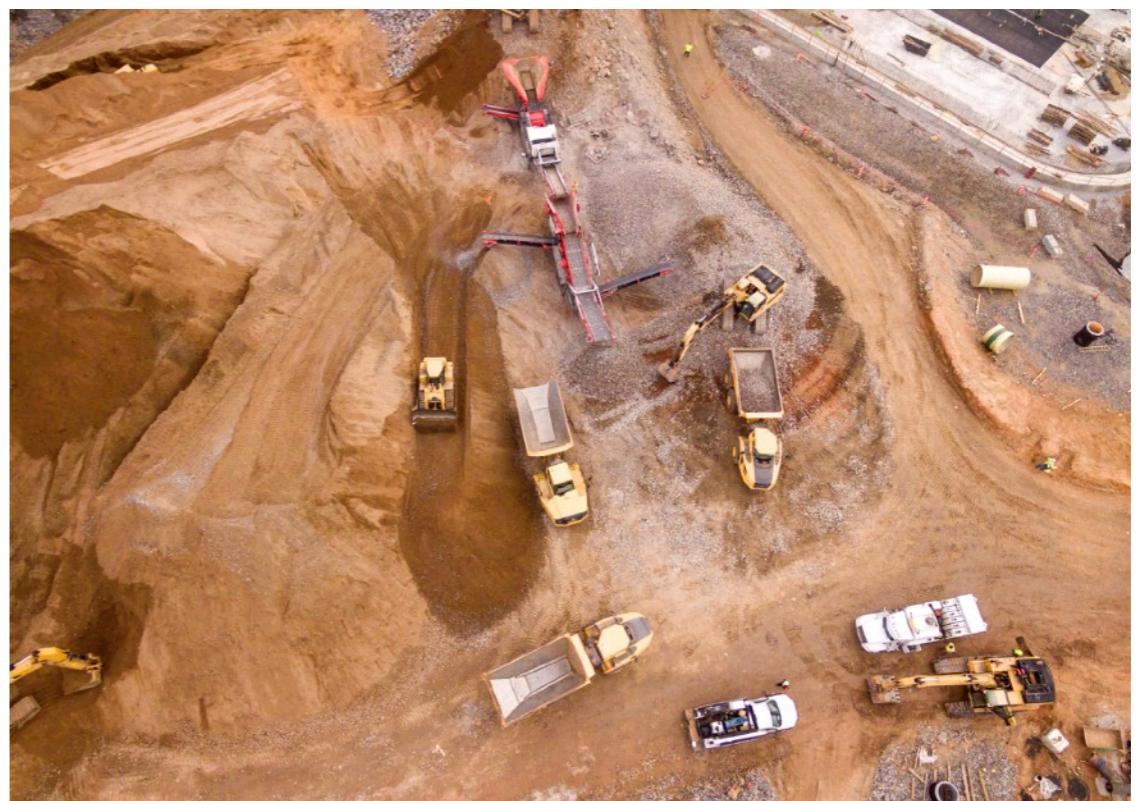
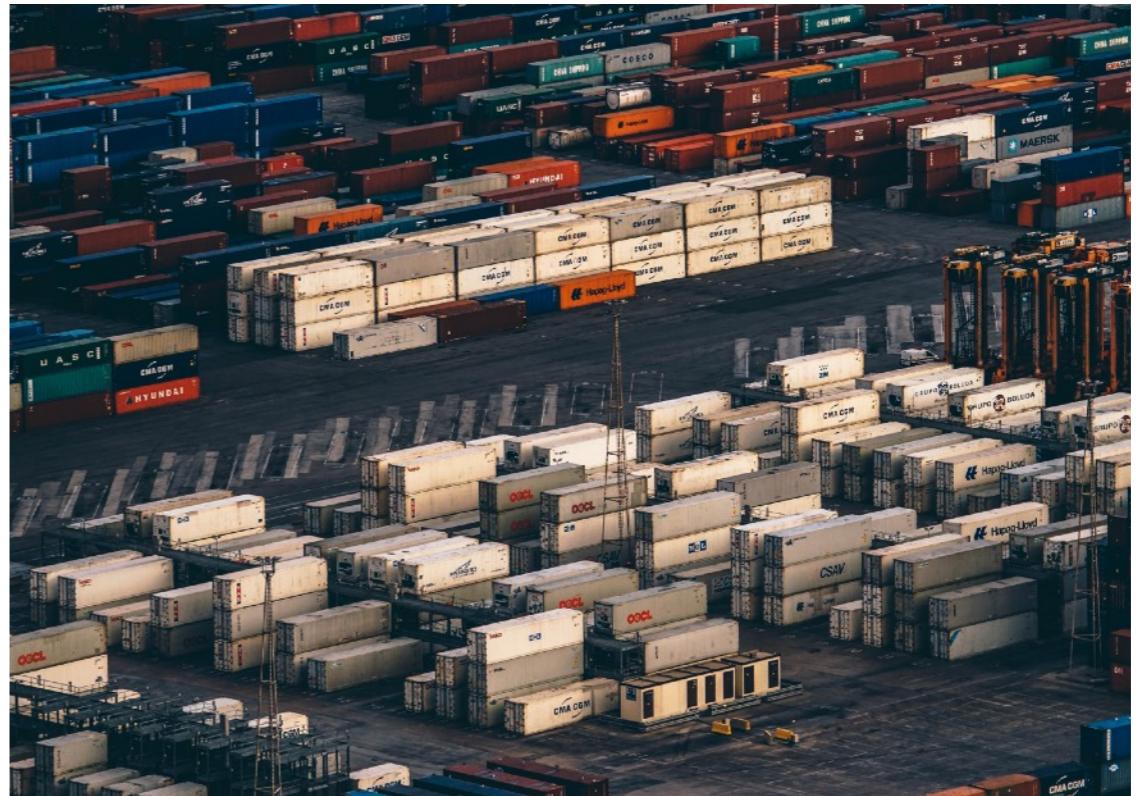


# Machine Learning for Combinatorial Optimization

Bengio, Lodi, Prouvost  
ElementAi  
21<sup>st</sup> of February 2019

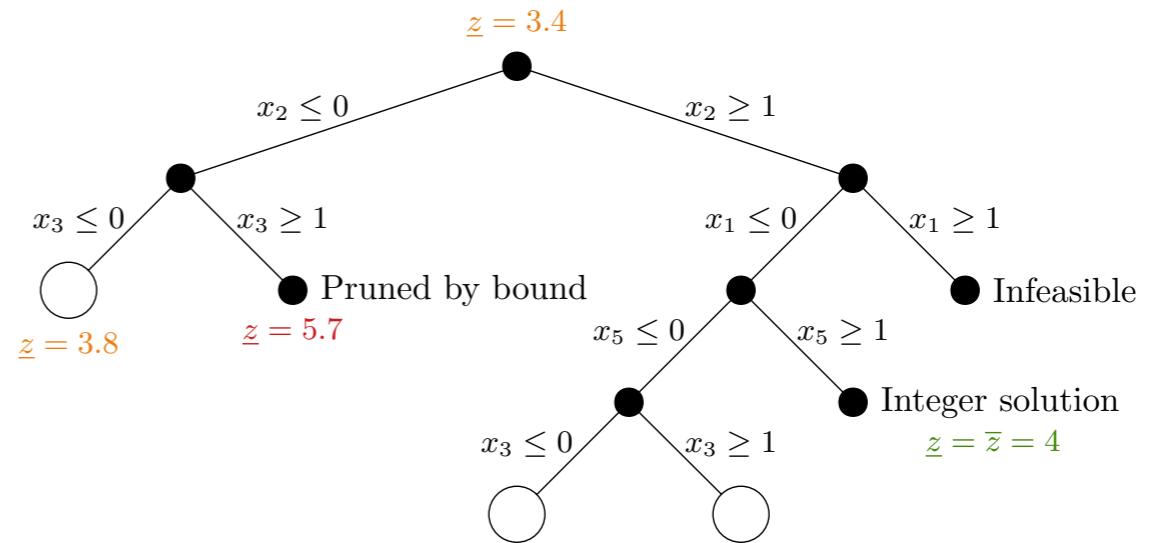






# Discrete decisions

- MILP is NP-complete
- Many different techniques
  - Branch and bound
  - Cutting planes
  - Decomposition
  - Heuristics
  - Presolving



*A Branch and Bound Tree*



## Too long

- Expert knowledge of how to make decisions
- Too expensive to compute
- Need for fast approximation

## Too heuristic

- No idea which strategy will perform better
- Need a well performing policy
- Need to discover policies

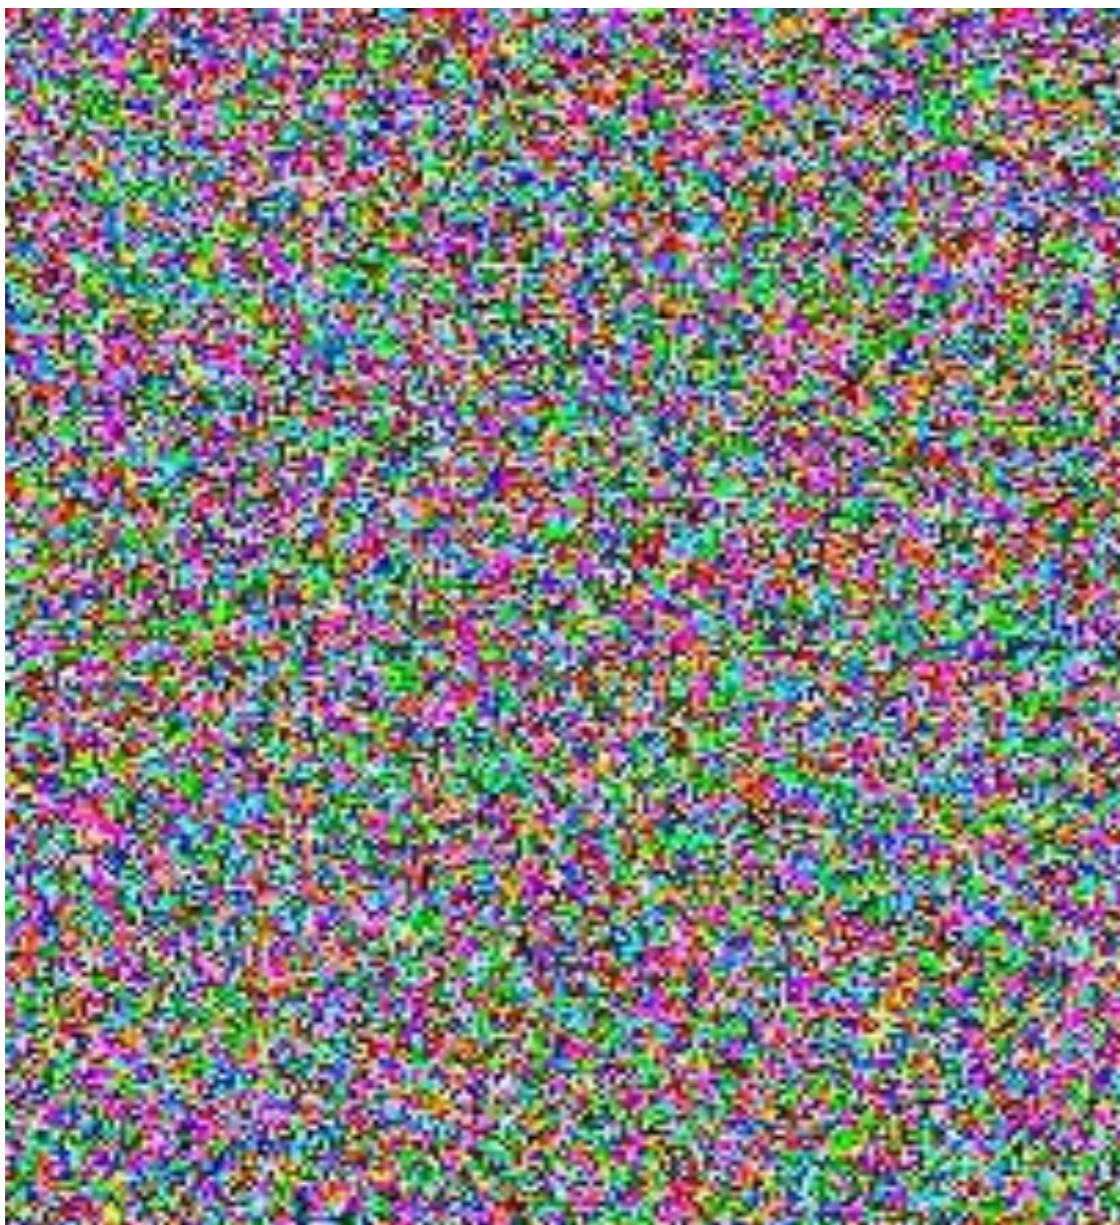
# Requirement

- We want to keep the **guarantees** provided by exact OR algorithms (feasibility, optimality)

# The structure hypothesis

- We do not care about most instances that could exist;
- Instead, we look at problem instances as data points from a specific, untractable, probability distribution;
- “Similar” instances show “similar” solving procedures.

# Random Images

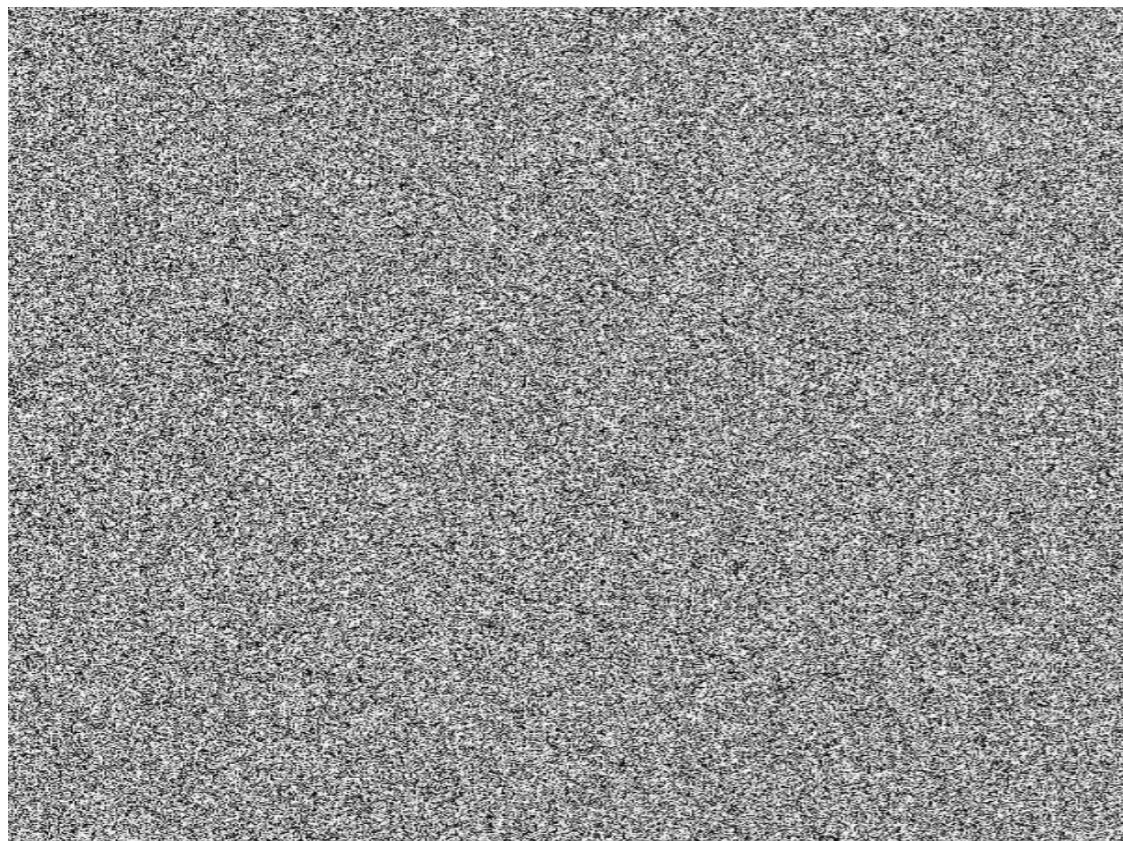


*Random iid pixels*

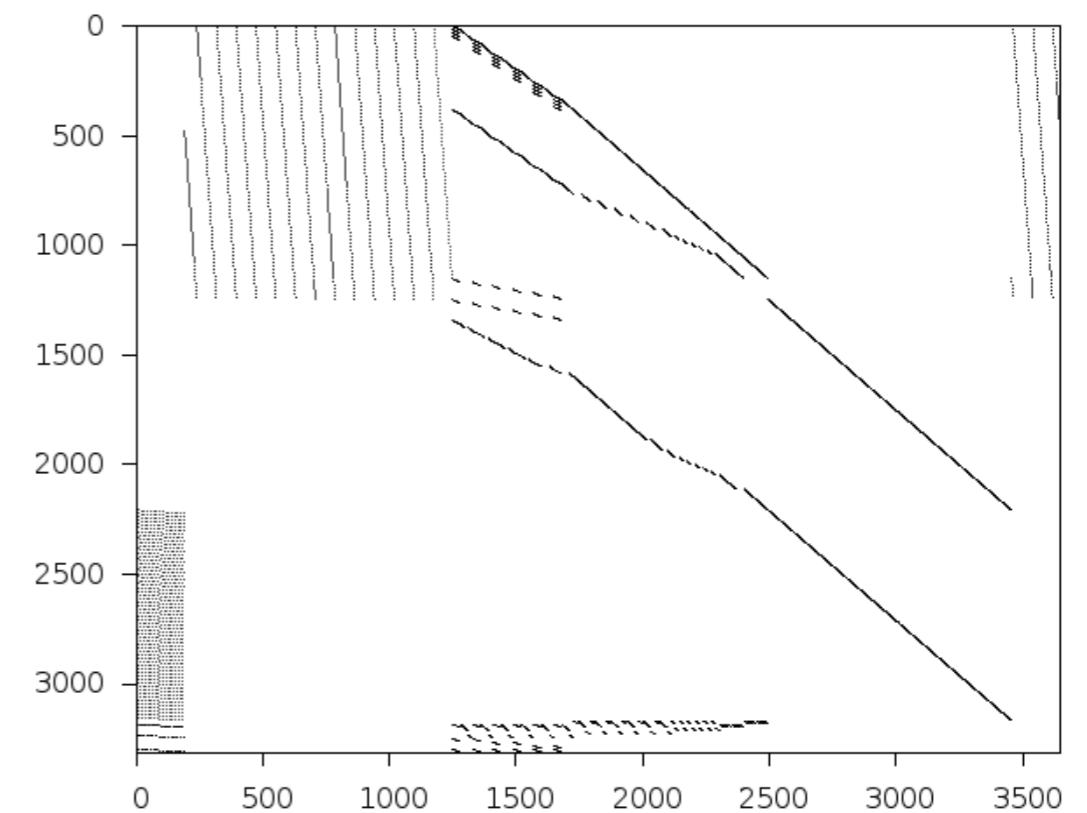


*Random face (GAN)*  
[thispersondoesnotexist.com](http://thispersondoesnotexist.com)

# Random Instances



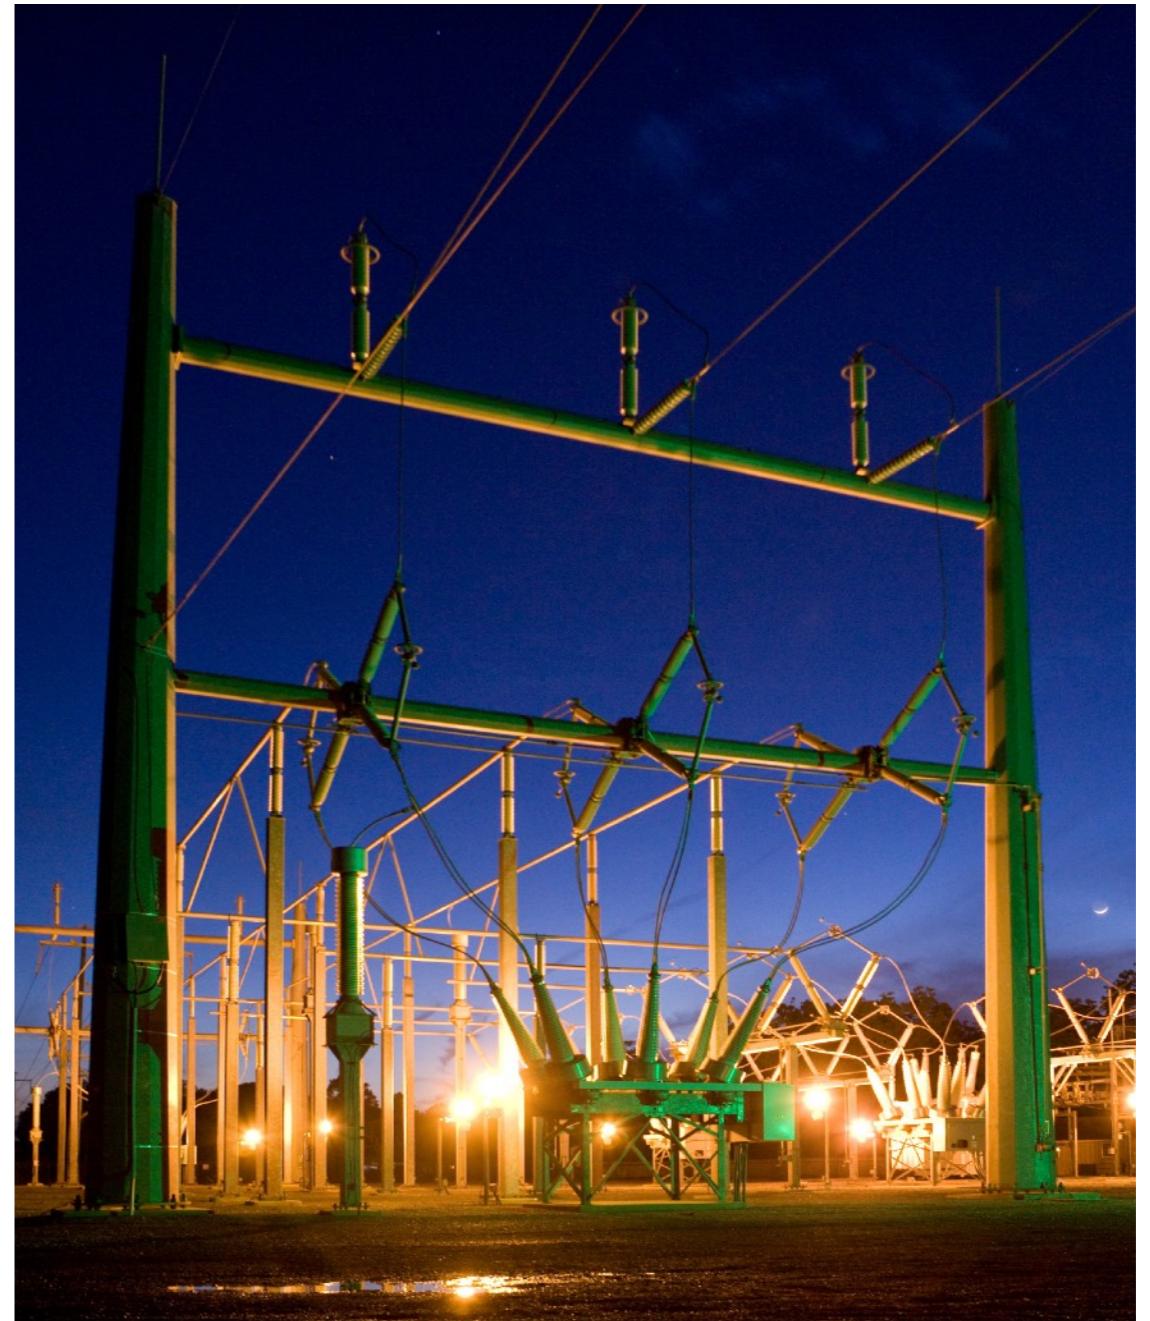
*Random iid coefficients*



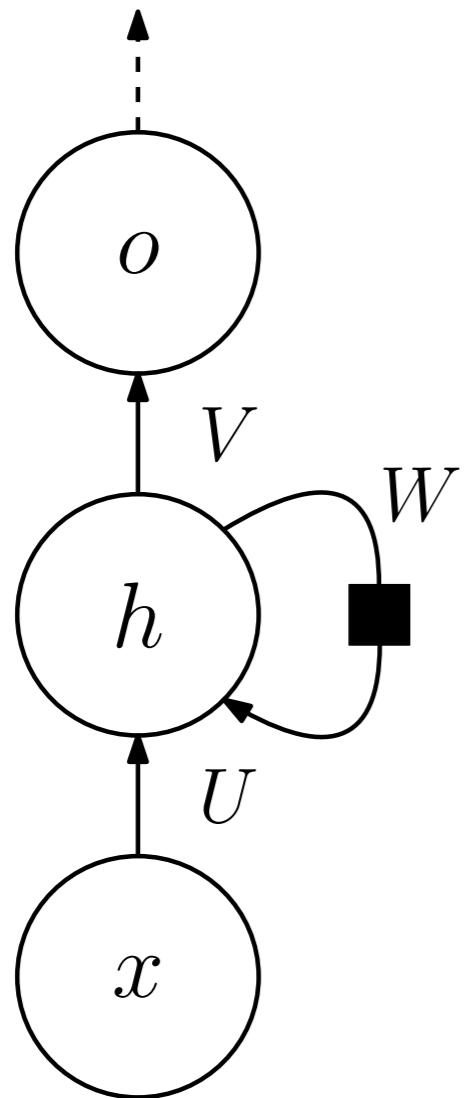
*a1c1s1 from MipLib 2017*

# Business Applications

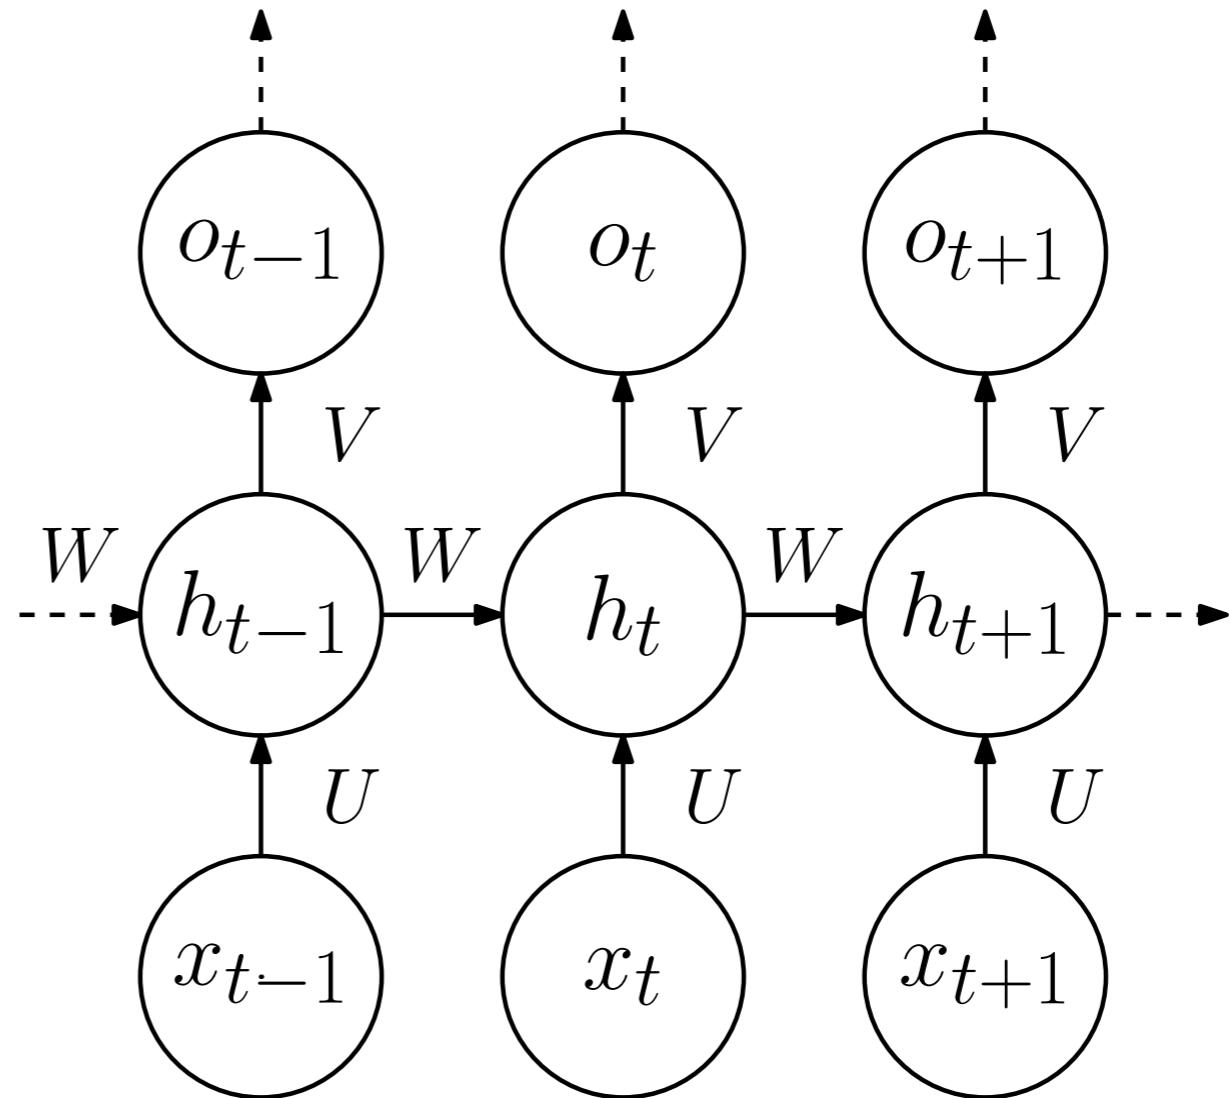
- Many businesses care about solving **similar** problems **repeatedly**
- Solvers do not make any use of this aspect
- Power systems and market  
[Xavier et al. 2019]
  - Schedule 3.8 kWh (\$400 billion) market annually in the US
  - Solved multiple times a day
  - 12x speed up combining ML and MILP



# Deep Learning Reminder

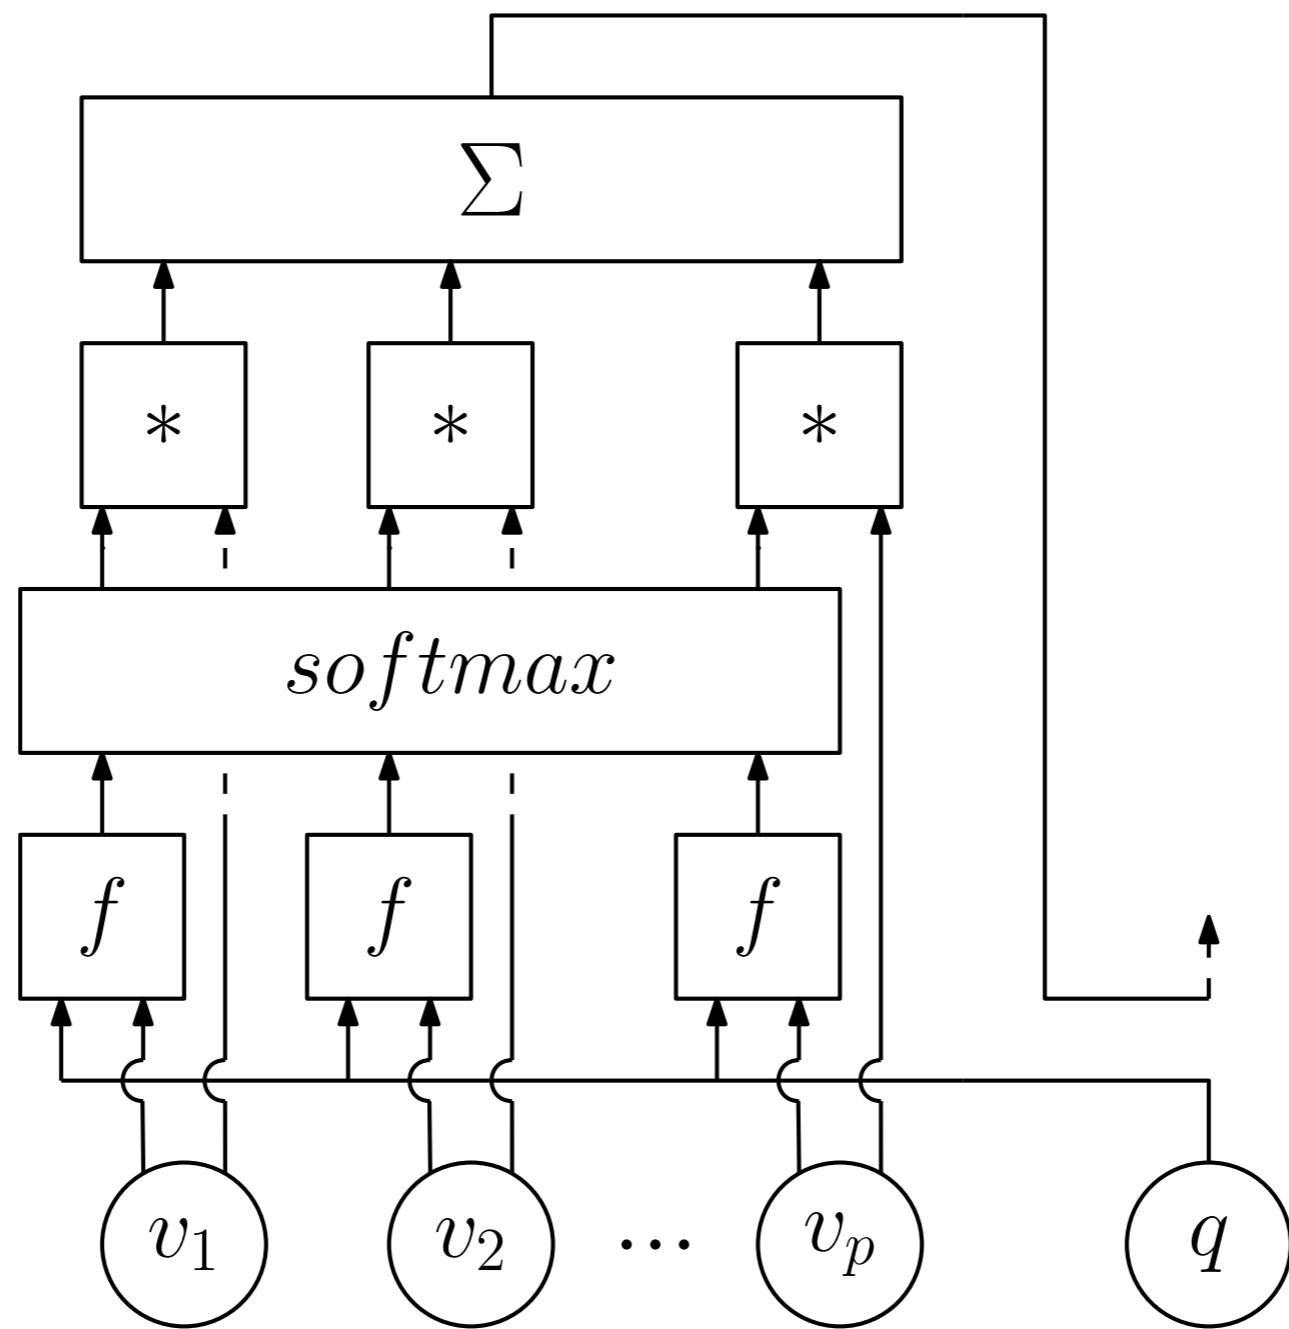


*RNN folded*



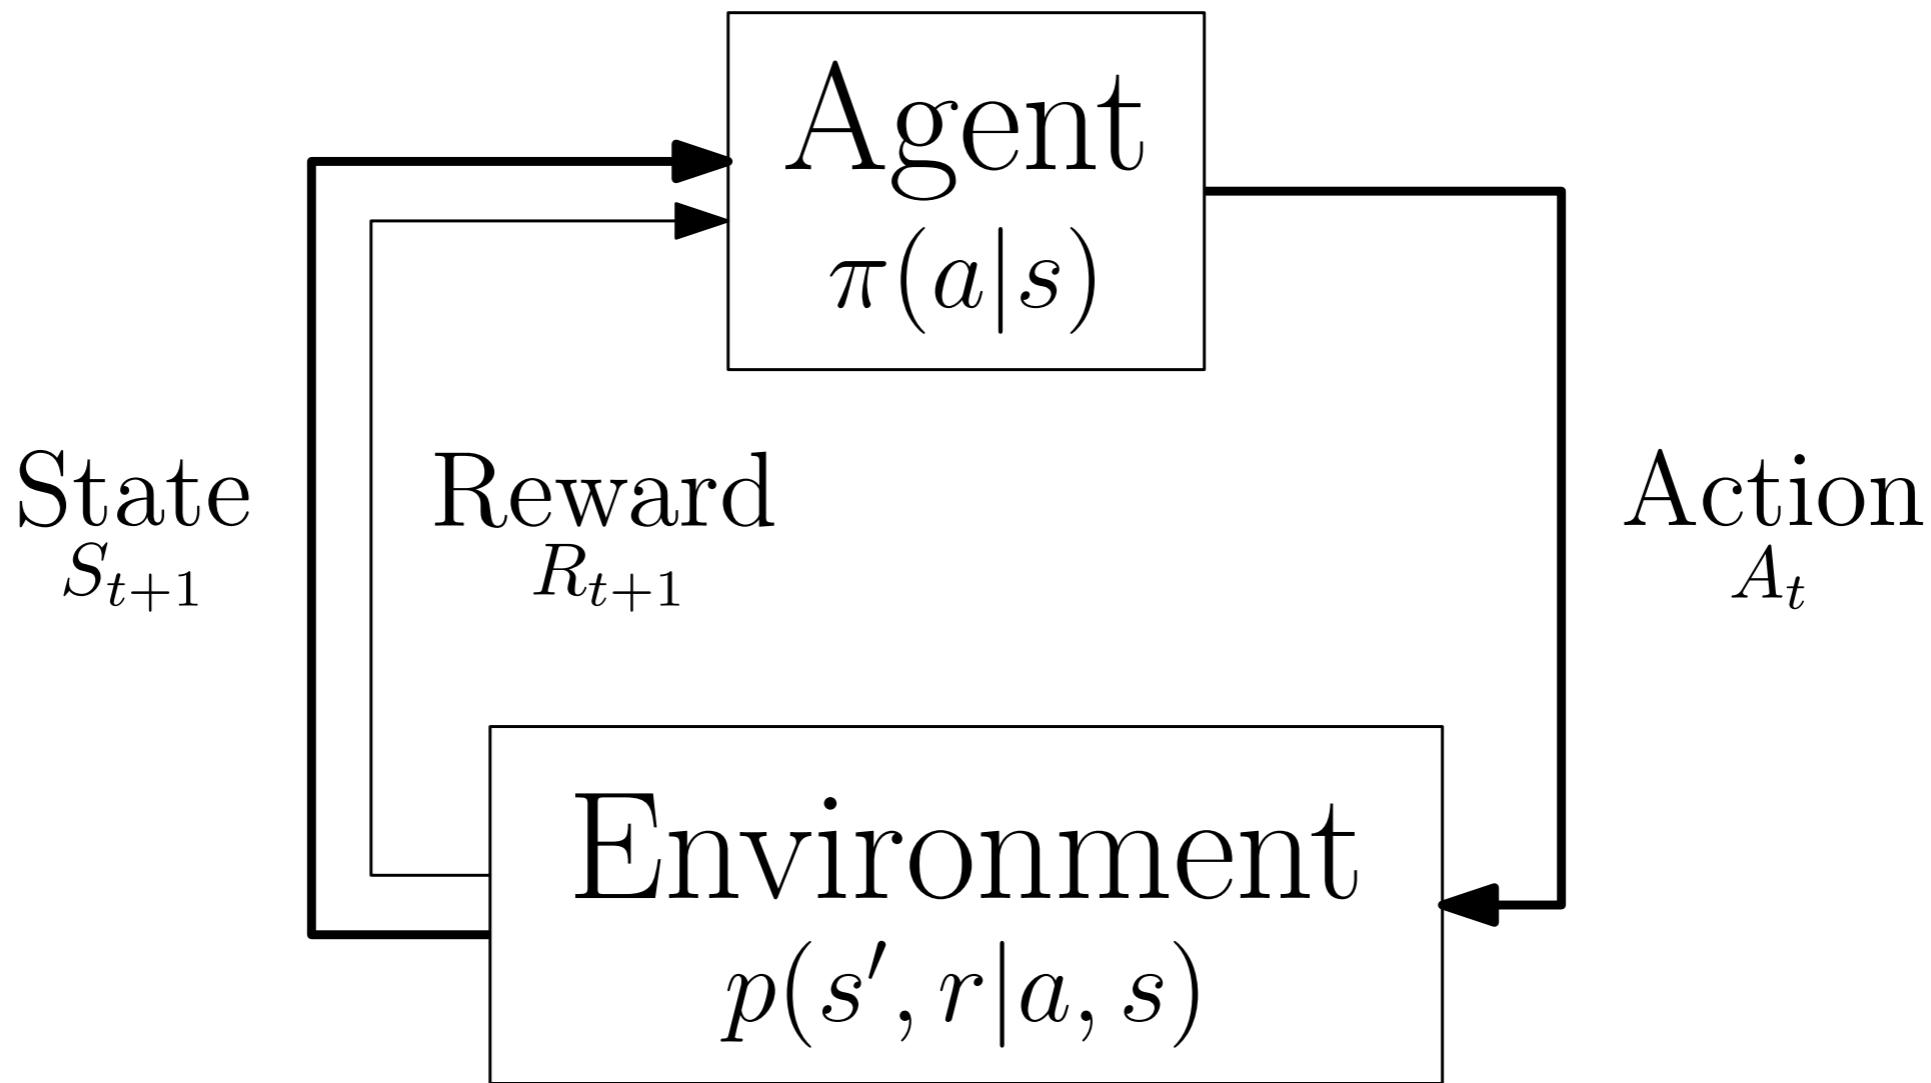
*RNN unfolded*

# Deep Learning Reminder



*Attention mechanism*

# Reinforcement Learning Reminder



*Markov Decision Process for Reinforcement Learning*

# Learning Methods

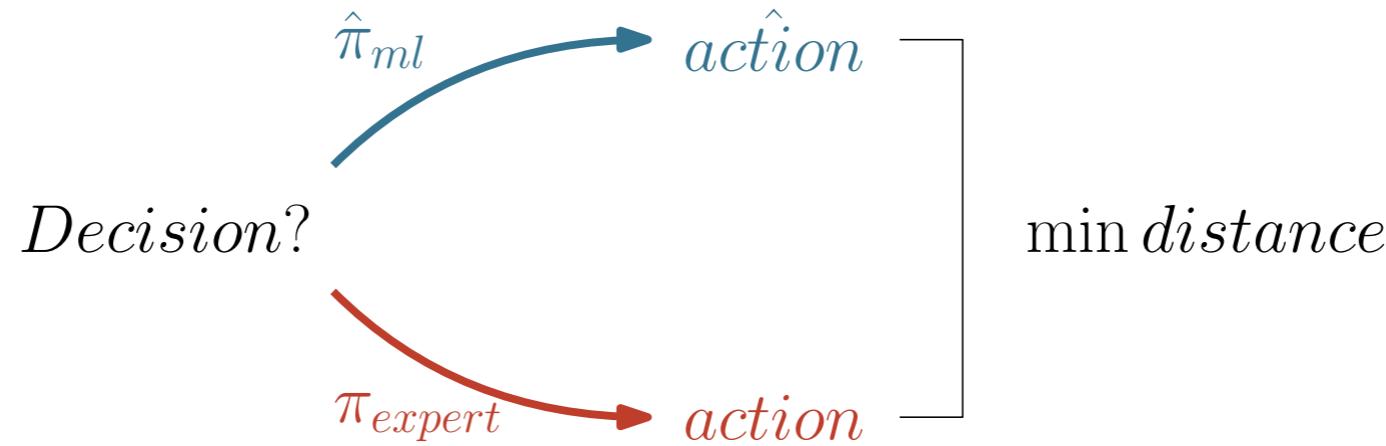
## Demonstration

- An expert/**algorithm** provides a policy
- Assumes theoretical / empirical knowledge about the decisions
- Decisions are too long to compute
- Supervised imitation learning

## Experience

- Learn and discover new policies (better hopefully)
- Unsatisfactory knowledge (not mathematically well-defined)
- Decisions are too heuristic
- Reinforcement learning

# Demonstration



- Approximating strong branching  
[Marcos Alvarez et al. 2014, 2016, 2017][Khalil et al. 2016]
- Approximating lower bound improvement for cut selection  
[Baltean-Lugojan et al. 2018]
- Approximating optimal node selection  
[He et al. 2014]

# Experience



- Learning greedy node selection (e.g. for TSP)  
[Khalil et al 2017a]
- Learning TSP solutions  
[Bello et al. 2017][Kool and Welling 2018][Emami and Ranka 2018]



## Not mutually exclusive

### **Supervised**

- Cannot beat the expert (an algorithm)  
→ only interesting if the approximation is faster
- Can be unstable
- Cannot cope with equally good actions

### **Reinforcement**

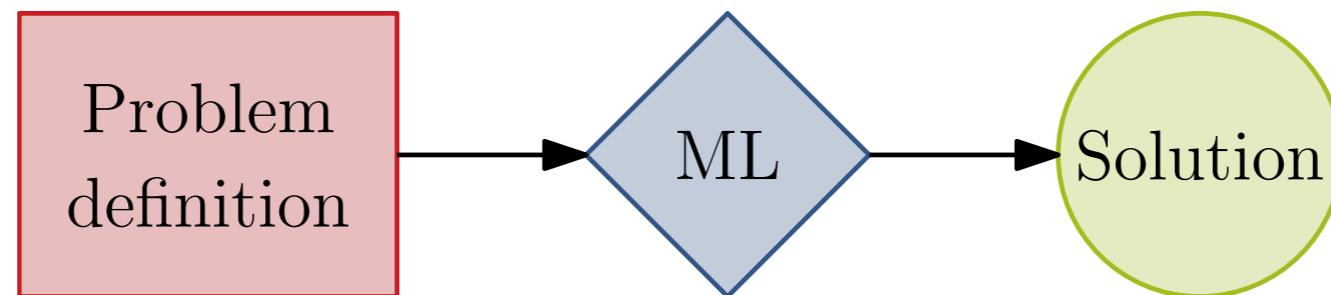
- Reinforcement can potentially discover better policies
- Harder, with local maxima (exploration difficult)
- Need to define a reward

**Better combined!**

# Algorithmic Structure

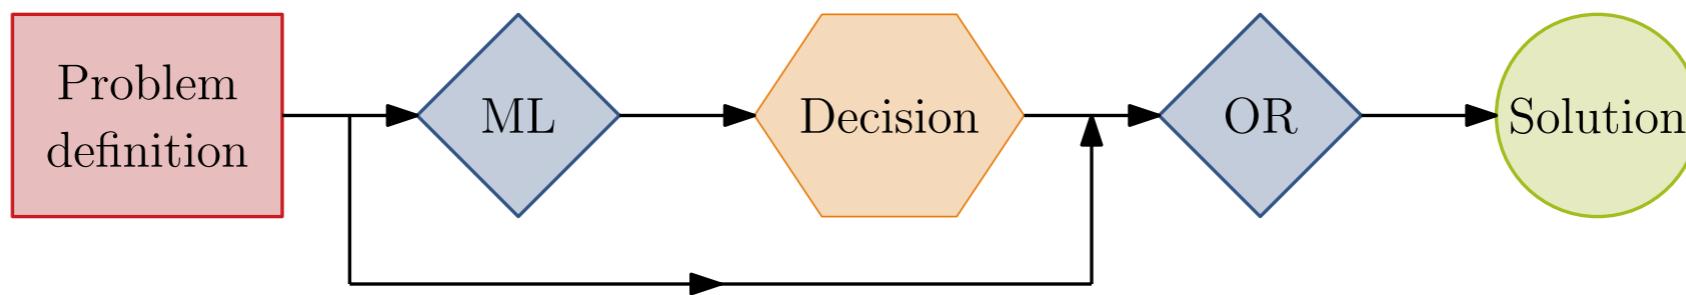
- *How do we build such algorithm? How do we mix OR with ML?*
- *How do we keep guarantees provided by OR algorithms (feasibility, optimality)?*

# End to End Learning



- Learning TSP solutions  
[Bello et al. 2017][Kool and Welling 2018][Emami and Ranka 2018]  
[Vinyals et al. 2015][Nowak et al. 2017]
- Predict aggregated solutions to MILP under partial information  
[Larsen et al. 2018]
- Approximate obj value to SDP (for cut selection)  
[Baltean-Lugojan et al. 2018]

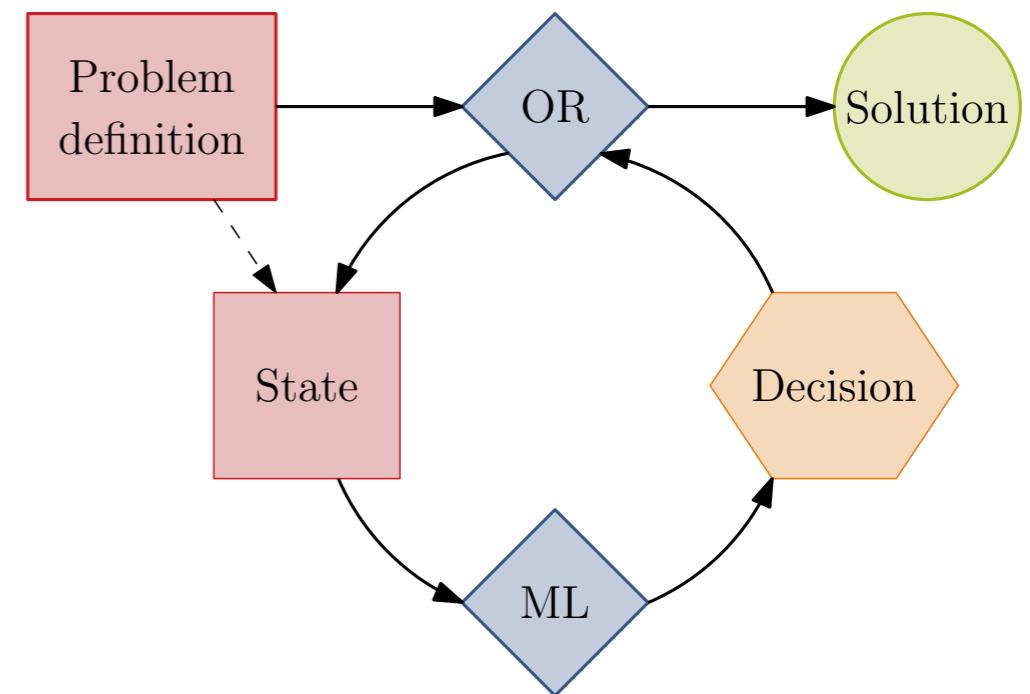
# Learning Properties



- Use a decomposition method  
[Kruber et al. 2017]
- Linearize an MIQP  
[Bonami et al. 2018]
- Provide initial cancer treatment plans to inverse optimization  
[Mahmood et al. 2018]

# Learning Repeated Decisions

- Learning where to run heuristics in B&B  
[Khalil et al. 2017b]
- Learning to branch  
[Lodi and Zarpellon 2017] (survey)
- Learning gradient descent  
e.g. [Andrychowicz et al. 2016]
- Predict booking decisions under unperfected information  
[Larsen et al. 2018]
- Learning cutting plane selection  
[Baltean-Lugojan et al. 2018]



Just a question  
of viewpoint

# Evaluation

- *What are our metrics?*
- *What instances do we want to generalize to?*
  - Instance specific policies should be easier to learn, but have to be re-learned every time
  - Policies that generalize can take some training offline (multi-task learning)
    - transfer learning, fine-tuning, meta-learning
- *What distribution of instances are we interested in?*

# Challenges

- *Which models and DL/RL algorithms will perform well?*
- *How do we represent the data? Should we approximate it?*
- *Can we scale?*
  - In the computations?
  - Generalizing?
  - Learning?
- *How to anticipate learning? Which distribution? How to generate data?*

# References and Paper

- Xavier, A. S., Qiu, F., & Ahmed, S. (2019). Learning to Solve Large-Scale Security-Constrained Unit Commitment Problems. ArXiv:1902.01697 [Cs, Math, Stat]. Retrieved from <http://arxiv.org/abs/1902.01697>
- All other references are listed in the paper:  
Bengio, Lodi, Prouvost (2018) - Machine Learning for Combinatorial Optimization: a Methodological Tour d'Horizon  
<https://arxiv.org/abs/1811.06128>